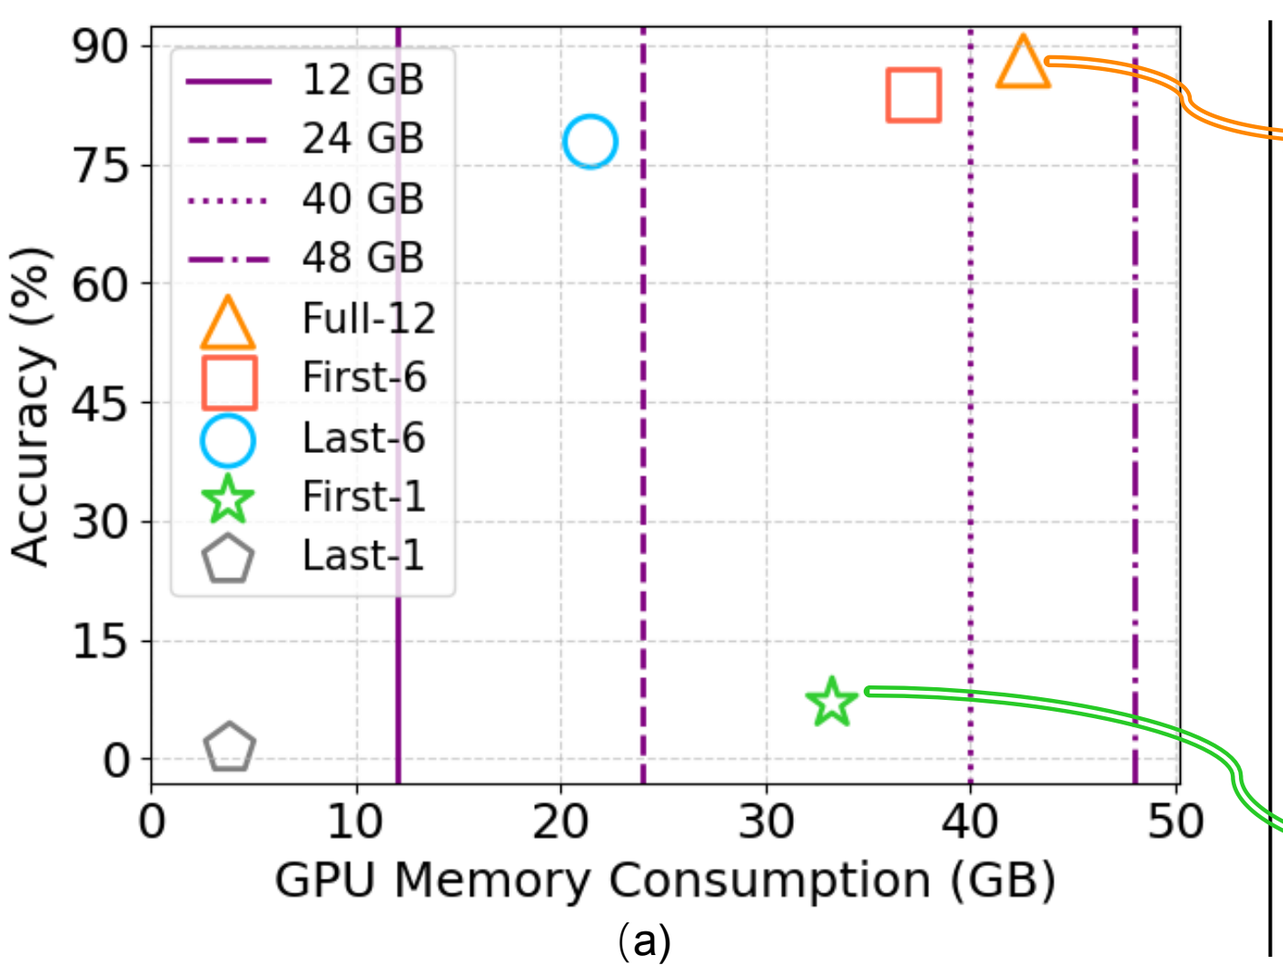
Accuracy (%)
GPU Memory Consumption (GB)
12 GB
24 GB
40 GB
48 GB
Full-12
First-6
Last-6
First-1
Last-1
0
15
30
45
60
75
90
0
10
20
30
40
50

(a)

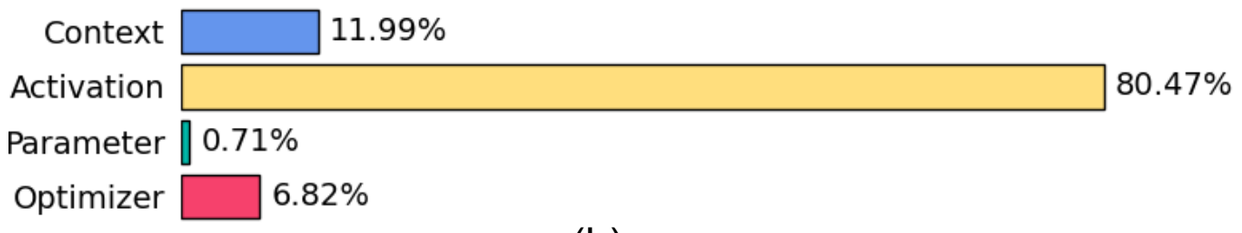
Context
11.99%
Activation
80.47%
Parameter
0.71%
Optimizer
6.82%

(b)

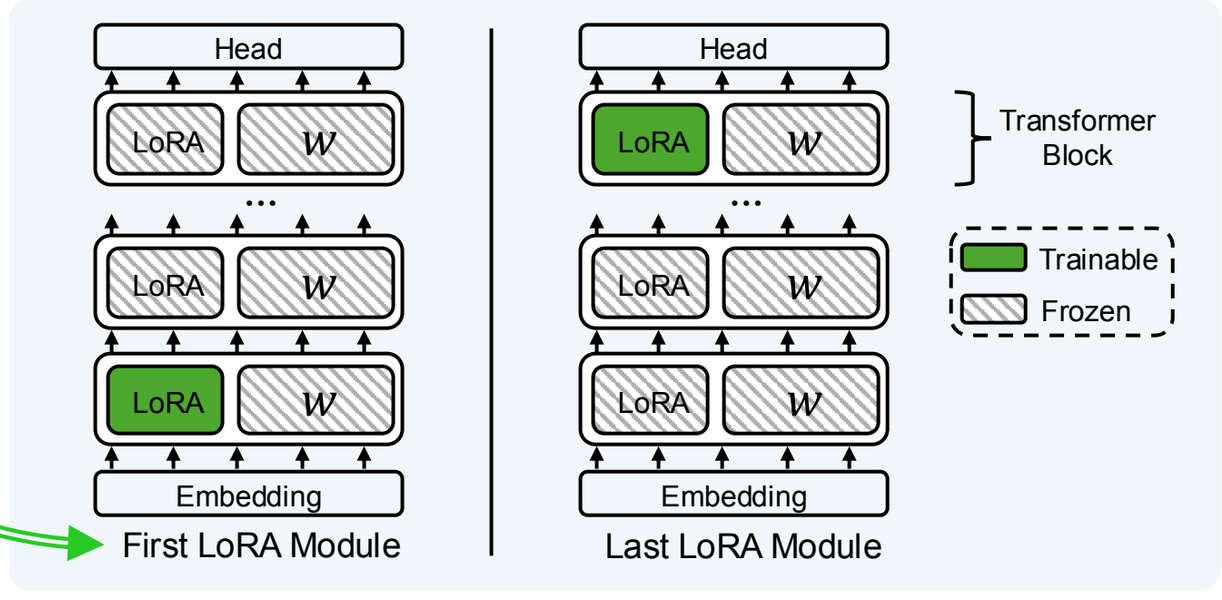
Head
LoRA
W
Embedding
First LoRA Module
Last LoRA Module
Transformer Block
Trainable
Frozen

(c)